BME 548L: Homework 2

February 24, 2021

Due date: Wednesday March 10, 2021 by 11:59pm
Please write up solutions to all problems (however you'd like) and submit them in class or via the Google Forms by the above due date. It is important that you show all of your steps – points will be deducted for simply writing the answer without showing any intermediate steps on how you arrived at your answer. You may work together to solve these problems, but please write up your own answer in your own way.

Remember that there is a coding component to this homework assignment as well, which can be found on the website: `https://deepimaging.github.io/homework/hw2`

Problem 1: The magic of the max operator. In this problem, we're going to try to classify a set of 3 different training data points in two dimensions. Here are the example points and their associated labels:

$$\mathbf{x}_1 = \begin{bmatrix} -1 \\ 0 \end{bmatrix}, y_1 = +1; \mathbf{x}_2 = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, y_2 = -1; \mathbf{x}_3 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, y_3 = +1$$

(a) Let's try to classify these points with a linear classifier, $y^* = \text{sign}(\mathbf{w}^T\mathbf{x})$, where $\mathbf{w}$ is the unknown 3×1 weight vector that we'd like to solve for. Remember, we set the first entry of each 3×1 data vector $\mathbf{x}$ to 1 to fold the usual constant offset term b into the inner product between $\mathbf{w}$ and $\mathbf{x}$. To solve for $\mathbf{w}$, please construct a 3×3 data matrix $\mathbf{X}$ that holds each of our 3 example points, expressed as a vector, in each matrix row. Then, solve for the optimal linear classifier weights $\mathbf{w}$ by computing the pseudo-inverse of $\mathbf{X}$, and then by multiplying this pseudo-inverse, $\mathbf{X}^\dagger$, with a 3×1 vector containing the labels. You can either compute $\mathbf{X}^\dagger$ and the resulting $\mathbf{w}$ by hand, or use a computer to help.

(b) How well do the optimal weights $\mathbf{w}$ do at classifying the original 3 data points? To check this, it is helpful to solve for $\mathbf{y}^* = \text{sign}(\mathbf{w}^T\hat{X})$, where $\hat{X}$ is a matrix that holds each of the training data points in its 3 *columns*, and $\mathbf{y}^* \in \mathbb{R}^{1\times3}$ is a vector that holds the three classification score estimates. Please sketch the 3 points in 2D space, and discuss why the above classification attempt was or was not so successful.

(c) Maybe it'll help to do a convolutional classification instead of a straight linear classification? To test this, let's assume that we've heard that a good convolutional filter to mix this data up a bit is $\mathbf{c} = [3, -1]$ (written as a row vector here to save space).

1

First, construct a convolution matrix $\mathbf{W}_1 \in \mathbb{R}^{4\times 3}$ with $\mathbf{c}$ in each column. Then, we'll try out the following classifier,

$$y^* = \text{sign}(\mathbf{w_2}^T \mathbf{W}_1 \mathbf{x}), \tag{1}$$

where $\mathbf{w_2}$ is a 4×1 weight vector of unknown weights. Is it possible to determine a $\mathbf{w_2}$ that can correctly classify the three points? Please provide a mathematical argument for why or why not.

(d) Finally, let's try to add a non-linearity in there, to see if that will help. Using the same $\mathbf{W}_1$ as above, let's solve,

$$\mathbf{y}^* = \text{sign}\left(\mathbf{w_2}^T \text{ReLU}(\mathbf{W}_1 \hat{\mathbf{X}})\right), \tag{2}$$

where $\mathbf{w_2}$ is again a 4×1 weight vector of unknown weights, and the ReLU operation is the max() operation that sets all values less than 0 equal to 0. Is it possible to determine a $\mathbf{w_2}$ that can accurately classify the three points contained in the columns of $\hat{\mathbf{X}}$? Please provide a mathematical argument for why or why not.

Problem 2: Cost function for logistic classification. We showed in class that the cross-entropy cost function works well for classifying images into outputs that are treated as probability distributions, as opposed to binary "yes/no" categories. This cost function takes the following form for the case of logistic classification:

$$L_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^{N} \ln\left(1 + e^{-y_n \mathbf{w}^T \mathbf{x_n}}\right) \tag{3}$$

(a) Noting that the standard logistic function $\theta(x) = e^x/1 + e^x$, please compute the gradient of L_{in} and show that,

$$\nabla_{\mathbf{w}} L_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^{N} -y_n \mathbf{x}_n \theta(-y_n \mathbf{w}^T \mathbf{x}_n) \tag{4}$$

(b) Based on this computation, please show that a misclassified input contributes more to the gradient than a correctly classified input. This is an important thing to remember - the inputs that are misclassified will be important in driving your gradient descent solver!

(c) (bonus problem) Compute the Hessian of $L_{in}(\mathbf{w})$ with respect to $\mathbf{w}$. Then, use this Hessian to write an expression for the optimal step size to use for a steepest descent solver for the logistic classification problem.

Problem 3: Perceptron Learning Algorithm. A simple iterative method related to solving the logistic classification problem is termed the Perceptron Learning Algorithm. In the simplest form of this approach, we'll start by guessing an initial separation boundary between training data points, where the data points are assigned known labels in one of of

two categories. The separation boundary is defined by a vector of weights $\mathbf{w}$, just like we've been considering in class. At iteration t, where $t = 1, 2, ...$, the weight vector will take on a value $\mathbf{w}(t)$. Then, the algorithm will pick one of the *currently misclassified* training data points $(\mathbf{x}(t), y(t))$ and will use it to update $\mathbf{w}(t)$. Since the example is misclassified, we have $y(t) \neq \text{sign}(\mathbf{w}^T(t)\mathbf{x}(t))$. The update rule is,

$$\mathbf{w}(t+1) = \mathbf{w}(t) + y(t)\mathbf{x}(t) \tag{5}$$

The PLA algorithm continues through the iterative lop until there are no longer any misclassified points in the training data set, and will eventually converge to a linear separator for linearly separable data.

(a) To show that the weight update rule above moves in the direction of correctly classifying $\mathbf{x}(t)$, first show that

$$y(t)\mathbf{w}^T(t)\mathbf{x}(t) < 0 \tag{6}$$

Hint: $\mathbf{x}(t)$ is misclassified by $\mathbf{w}(t)$

(b) Then, show that,

$$y(t)\mathbf{w}^T(t+1)\mathbf{x}(t) > y(t)\mathbf{w}^T(t)\mathbf{x}(t) \tag{7}$$

Hint: Use the update rule here

(c) As far as classifying $\mathbf{x}(t)$ is concerned, argue that that move from $\mathbf{w}(t)$ to $\mathbf{w}(t+1)$ is a move in the right direction. Feel free to draw a picture of a simple example in 2D if you find that helpful.

Problem 4: Let's write out a CNN equation!. Consider the CNN network for image classification sketched on the next page. We're going to write this out as a series of matrix operations.

Fig. 1(a) contains a simple 3-layer CNN network for 2D images, where a first convolution layer is comprised of 2 convolutional filters with unknown weights (which creates 2 depth slices). After this convolution, a nonlinearity is followed by "sum-pooling", which is an operation that takes every group of 2×2 pixels in the input and adds them together to form 1 new pixel at the output, effectively shrinking each image by a factor of 2 in either dimension. These 3 steps are repeated in Layer 2, where now the smaller set of 2 images are each convolved with a set of 2 convolutional filters with unknown weights to produce 4 depth slices. The final fully connected layer maps all of the data to a set of 3 possible categories via a full matrix multiplication. Since it would be difficult to write out full matrix equations for this pipeline, we consider a simpler case.

In Fig. 1(b), I have sketched the same pipeline for a 1D image, to make things a little simpler. It follows the same network architecture, but only deals with 1D data. Let's assume the 1D input image contains 4 pixels and is defined by: $\mathbf{x}_n \in \mathbb{R}^{1\times 4}$. Please write out this entire CNN pipeline as a set of matrix operations, where you show the correct size and composition of the elements within each matrix: 2 convolution matrices, 2 sum-pooling matrices and one fully connected matrix (i.e., a "Dense" layer). The final expression should contain a series of operations that are cascaded together to mathematically link the input to the output. Here are some notes to help you do that:

(a) CNN for 2D images - an example:

(b) CNN for a 1D image – please write out matrix operations for the following:

Figure 1: CNN layout for Problem 4

- You can just write 'ReLU' where a nonlinearity is applied, or ignore them in your analysis (you do not and cannot easily write these nonlinearities as matrix operations).

- I'd like to see all the matrices written out to establish the mapping between $\mathbf{x}_n \in \mathbb{R}^{1 \times 4}$ and $\mathbf{y}_n \in \mathbb{R}^{1 \times 3}$.

- You should express the convolutional filters as variables within each matrix, $\mathbf{c}_k^l = (c(1)_k^l, c(2)_k^l)$ for a 2x1 filter, using k to denote the depth slice and l the layer it applies to.

- To deal with the creation of multiple depth slices, you should "stack" matrices on top of one another.

- If needed, you can "pad" the input, as discussed in class